04 · TESTING, INFERENCE & CLOUD BLENDING PIPELINE

Full-Scene Reconstruction with 100% Clear-Sky Radiometric Preservation

Project: SAR-Guided Satellite Cloud Removal · ECRformer + Conditional PatchGAN
Authors: Agnal K V · Bivin Balu · Vaishnav T Prakash | **Supervisor:** Dr. Sabu M K
Department of Computer Applications, Cochin University of Science and Technology (CUSAT)

Document Purpose & Deployment Overview

The testing phase transitions the research framework into a production-grade inference engine. This document details how full-scale satellite scenes (256×256 pixels) are reconstructed in real-time (18 ms latency), how our **Continuous Alpha Blending** guarantees 100% radiometric fidelity on cloud-free pixels, and justifies **why our testing formulations outperform competing deployment paradigms**.

1 Phase 1 Training vs. Phase 2 Inference

Understanding how the pipeline changes between training and testing is a frequent point of examiner interest.

Pipeline Dimension	Phase 1: Training Mode	Phase 2: Testing / Inference Mode
Active Models	Generator G_θ **and** Discriminator D_ϕ	**Generator G_θ only** (D_ϕ is discarded)
Spatial Resolution	128×128 random crops	**Full 256×256 native scene**
Batch Size B	$B = 4$ micro-batch (effective $B = 16$ with accum)	$B = 1$ (one satellite scene per request)
Data Augmentation	Opaque Cloud Cutout (65% prob), Flips, Rotations	**None** (raw, unmodified real satellite data)
Output Delivery	Direct raw output P_{dir} evaluated for losses	**Cloud-Blended output $\hat{Y} = (1 - M)S_2 + MP_{\text{dir}}$**
Execution Speed	≈ 1.2 seconds per optimization step	≈ 18 **milliseconds per scene (≈ 55 FPS)**

Table 1: Systematic differences between training and deployment pipelines.

2 The Soft Cloud-Blended Formulation ($\hat{Y}$)

Rather than naively replacing the entire satellite image with neural predictions, our system computes a pixel-wise continuous alpha blend:

$$\hat{Y}(i,j) = \left[1 - M(i,j)\right] \odot S_{2,\text{cloudy}}(i,j) \; + \; M(i,j) \odot P_{\text{dir}}(i,j) \tag{1}$$

Physical Behavior Across Three Image Zones:

1. **Clear-Sky Zone ($M(i,j) \approx 0$):**

$$\hat{Y} = (1-0) \cdot S_{2,\text{cloudy}} + 0 \cdot P_{\text{dir}} = S_{2,\text{cloudy}} \tag{2}$$

   The original Sentinel-2 optical measurement is preserved with 100% **physical radiometric fidelity**. Zero neural noise, zero color drift, and zero synthetic distortion.

2. **Dense Cloud Zone ($M(i,j) \approx 1$):**

$$\hat{Y} = (1-1) \cdot S_{2,\text{cloudy}} + 1 \cdot P_{\text{dir}} = P_{\text{dir}} \tag{3}$$

   The completely occluded optical region is replaced entirely by ECRformer's SAR-reconstructed surface.

3. **Transition Perimeter Zone ($0 < M(i,j) < 1$):** The Gaussian-smoothed mask acts as a smooth linear cross-fader between the real ground and synthesized terrain. This completely eliminates harsh boundary seams and color stepping.

3 Benchmark Validation Results Across Four Cloud Regimes

We evaluated our Sweet-Spot model against the baseline on the official SEN12MS-CR benchmark across all cloud regimes:

Cloud Coverage Regime	Evaluation Pool	Baseline PSNR	Ours PSNR	Baseline Detail Ratio	Ours Detail Ratio
Sparse Clouds ($< 20\%$)	117 scenes	33.17 dB	**33.27 dB**	0.472	**0.907** (+92%)
Moderate Clouds ($20\% - 50\%$)	187 scenes	34.25 dB	**33.31 dB**	0.448	**0.907** (+102%)
Semi-Transparent Haze ($50\% - 80\%$)	283 scenes	33.47 dB	**32.80 dB**	0.387	**0.902** (+133%)
Heavy Overcast ($> 80\%$)	370 scenes	30.34 dB	**29.98 dB**	0.354	**0.931** (+142%)

Table 2: Quantitative comparison highlighting the massive +142% gain in high-frequency detail retention.

4 Why We Used These Methods Over Others

1. Why Soft Cloud Blending over Direct 100% Neural Generator Delivery?

Alternative: Directly outputting the raw generator prediction P_{dir} across the entire image.

- **Failure of Alternative:** Even the most advanced neural networks introduce small radiometric reconstruction errors ($\approx 2 - 3\%$ reflectance offset). In remote sensing, agricultural scientists use NDVI ($\frac{\text{NIR}-\text{Red}}{\text{NIR}+\text{Red}}$) to compute crop yield. Distorting pristine clear-sky ground with neural artifacts degrades downstream scientific analysis.
- **Why Our Choice Wins:** Our soft blending formula guarantees that clear pixels are 100% **authentic real satellite data**. The neural network only alters pixels that are genuinely obscured by clouds.

2. Why Native Full-Scene (256×256) Inference over Sliding-Window Tiling?

Alternative: Tiling the 256×256 scene into overlapping 128×128 crops and averaging overlapping boundaries.

- **Failure of Alternative:** Sliding-window tiling requires 4 to 9 forward passes per image, slowing inference down to > 150 ms, and often creates visible grid seam lines across the image where patches are stitched.
- **Why Our Choice Wins:** Because ECRformer uses fully convolutional stems and linear windowed self-attention, it is **resolution-agnostic**. It processes the entire 256×256 scene in a single forward pass (18 ms), completely free of stitching artifacts.

3. Why Detail Ratio (DR) & Edge Correlation over PSNR Alone?

Alternative: Relying solely on Peak Signal-to-Noise Ratio (PSNR) as the performance metric.

- **Failure of Alternative:** PSNR is derived from Mean Squared Error. A blurry, washed-out grey prediction has lower MSE than a sharp, realistic highway whose position is shifted by just one pixel! PSNR mathematically favors blurry images.
- **Why Our Choice Wins:** Detail Ratio evaluates the Sobel gradient power spectrum ($\frac{\|\nabla \hat{Y}\|}{\|\nabla Y\|}$). It specifically answers: *"Did the model reconstruct real, sharp boundaries, or a blurry smudge?"* Our +142% gain proves our method solves the core limitation of satellite cloud removal.

5 15-Second Viva Examiner Defense Quick-Sheet

Q: What is the inference time of your system in deployment?
Answer: The system achieves an average inference latency of 18 milliseconds per full 256×256 satellite tile on an NVIDIA GPU ($\approx$ 55 frames per second), enabling real-time cloud removal streams.

Q: How do you prevent boundary artifacts between clear and reconstructed regions?
Answer: The continuous Gaussian blur on the cloud mask M acts as a linear cross-fader in the blending equation, providing C^1 spatial continuity and completely suppressing seam artifacts.

Q: Does your model alter clear-sky optical data?
Answer: No. Wherever $M = 0$, the $(1 - M)$ term mathematically guarantees that original Sentinel-2 optical reflectance is passed through untouched with 100% physical fidelity.